

```
In [1]: import numpy as np
import pandas as pd
import scipy.stats as stats
import statsmodels.api as sm
import statsmodels.formula.api as smf
from statsmodels.formula.api import ols
import matplotlib.pyplot as plt
import math
from scipy.stats import norm
import seaborn as sns
```

```
In [2]: # 1a
battery = np.array([25.5, 26.1, 26.8, 23.2, 24.2, 28.4, 25.0, 27.8, 27.3, 25.7])
mean_batt = np.mean(battery)
std_batt = np.std(battery, ddof=1)
n = len(battery)
```

```
In [3]: t_stat, p_val_two = stats.ttest_1samp(battery, 25)
p_val_one = p_val_two / 2 if t_stat > 0 else 1 - p_val_two/2
print(f"sample mean = {mean_batt:.3f}")
print(f"std = {std_batt:.3f}")
print(f"n = {n}")
print(f"t = {t_stat:.3f}")
print(f"one-sided p = {p_val_one:.4f}")
```

```
sample mean = 26.000
std = 1.625
n = 10
t = 1.946
one-sided p = 0.0417
```

```
In [4]: print(f" Conclusion that 5% significance level demonstrates population mean is greater than the hypothesized me:
Conclusion that 5% significance level demonstrates population mean is greater than the hypothesized mean.
```

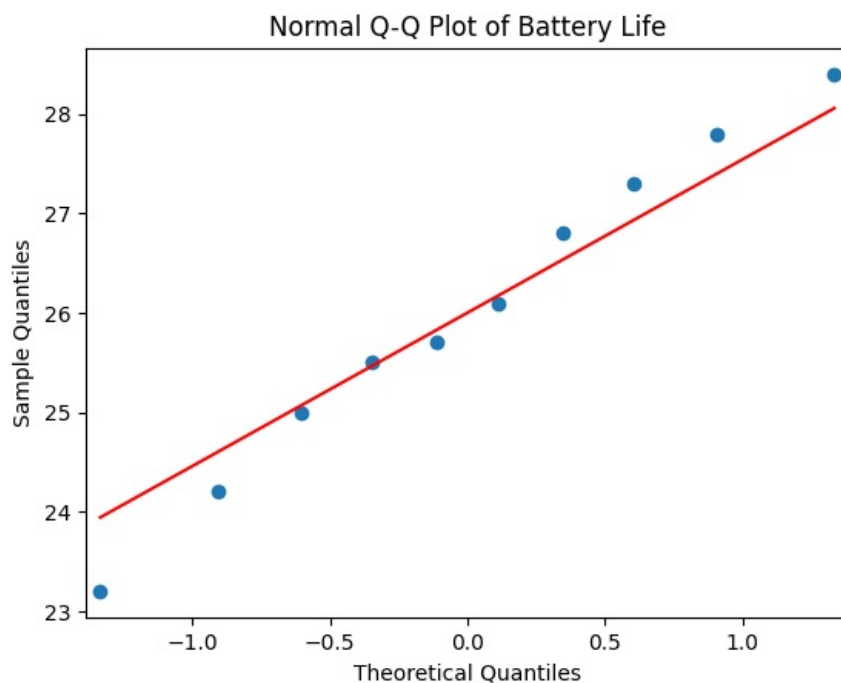
```
In [5]: alpha = 0.10
t_crit = stats.t.ppf(1 - alpha/2, n - 1)
margin = t_crit * std_batt / np.sqrt(n)
ci_lower, ci_upper = mean_batt - margin, mean_batt + margin
print(f"90% CI for mean:")
print(f"lower: {ci_lower:.3f}")
print(f"upper: {ci_upper:.3f}")
```

```
90% CI for mean:
lower: 25.058
upper: 26.942
```

```
In [6]: print(f"Tells me hypothesis confirmed, population mean is larger than 25.")
```

```
Tells me hypothesis confirmed, population mean is larger than 25.
```

```
In [7]: # 1c
sm.qqplot(battery, line='s')
plt.title("Normal Q-Q Plot of Battery Life")
plt.show()
```



```
In [8]: print(f"Conclusion is assumption of normality. Data points along approximately straight line.")
```

Conclusion is assumption of normality. Data points along approximately straight line.

```
In [9]: # 2a
n = 500
x = 65
p_hat = x / n
p0 = 0.08
alpha = 0.10
```

```
In [10]: se0 = math.sqrt(p0 * (1 - p0) / n)
```

```
In [11]: z = (p_hat - p0) / se0
p_value = 2 * (1 - norm.cdf(abs(z)))
```

```
In [12]: print(f" z statistic = {z:.3f}")
print(f" two-sided p-value = {p_value:.5f}")

z statistic = 4.121
two-sided p-value = 0.00004
```

```
In [13]: if p_value < alpha:
    print("Reject H0: p ≠ 0.08")
else:
    print("Fail to reject H0: p ≠ 0.08")
```

Reject H0: p ≠ 0.08

```
In [14]: # 2b
z_95 = norm.ppf(0.95) # critical value for 95% one-sided
se_hat = math.sqrt(p_hat * (1 - p_hat) / n)
upper_bound = p_hat + z_95 * se_hat
```

```
In [15]: print(f"Upper bound = {upper_bound:.3f}")
print(f"Interpretation is 95% confident that true process fraction nonconforming is no greater than {upper_bound:.3f}")
```

Upper bound = 0.155

Interpretation is 95% confident that true process fraction nonconforming is no greater than 15.5%.

```
In [16]: # 3
micrometer = np.array([0.150,0.151,0.151,0.152,0.151,0.150,0.151,0.153,0.152,0.151,0.151,0.151])
vernier = np.array([0.151,0.150,0.151,0.150,0.151,0.151,0.153,0.155,0.154,0.151,0.150,0.152])
diff = micrometer - vernier
t_stat3, p_val3 = stats.ttest_rel(micrometer, vernier)
print(f"There is a mean difference.")
print(f"Mean difference = {np.mean(diff):.6f}")
```

There is a mean difference.

Mean difference = -0.000417

```
In [17]: # 4a
flow = [125,125,125,125,125,125,
        160,160,160,160,160,160,
        200,200,200,200,200,200]
```

```
In [18]: uniformity = [2.7,2.6,4.6,3.2,3.0,3.8,
                      4.6,4.9,5.0,4.2,3.6,4.2,
                      4.6,2.9,3.4,3.5,4.1,5.1]
```

```
In [19]: df = pd.DataFrame({'Flow': flow, 'Uniformity': uniformity})
```

```
In [20]: model = ols('Uniformity ~ C(Flow)', data=df).fit()
anova_table = sm.stats.anova_lm(model, typ=2)
```

```
In [21]: print("\nANOVA Table:")
print(anova_table)
```

ANOVA Table:

	sum_sq	df	F	PR(>F)
C(Flow)	3.647778	2.0	3.585627	0.053365
Residual	7.630000	15.0	NaN	NaN

```
In [22]: alpha = 0.05
p_value = anova_table["PR(>F)"][0]
```

```
/tmp/ipykernel_966389/2957397728.py:2: FutureWarning: Series.__getitem__ treating keys as positions is deprecated. In a future version, integer keys will always be treated as labels (consistent with DataFrame behavior). To access a value by position, use `ser.iloc[pos]`
p_value = anova_table["PR(>F)"][0]
```

```
In [23]: if p_value < alpha:
    print(f"p = {p_value:.4f} < 0.05")
```

```

    print(f"Reject H0. Flow rate significantly affects etch uniformity.")
else:
    print(f"p = {p_value:.4f} ≥ 0.05")
    print(f"Fail to reject H0. Flow rate does not significantly affect etch uniformity.")

```

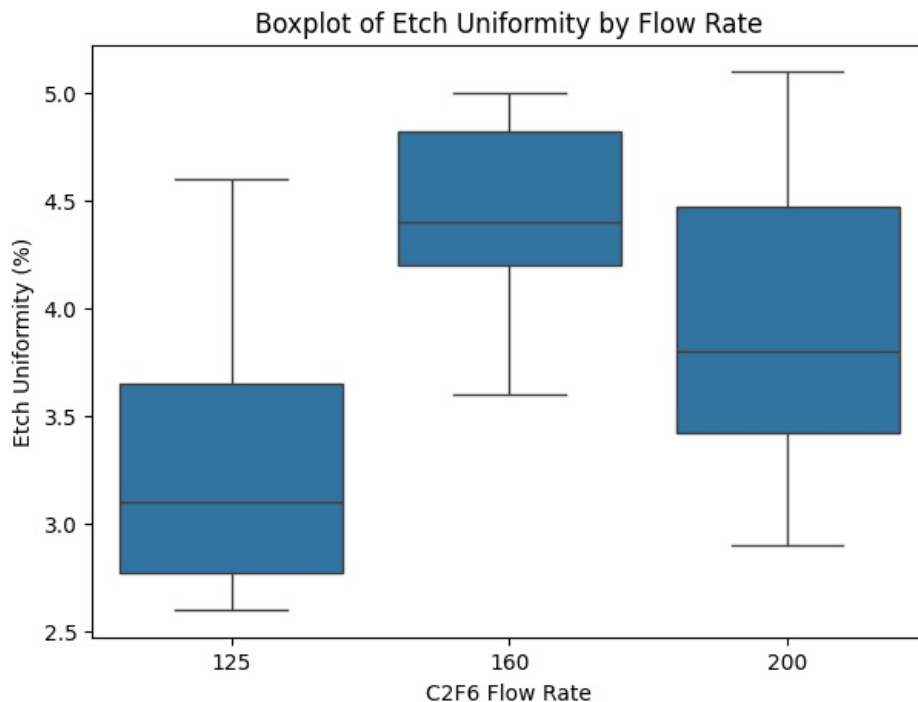
p = 0.0534 ≥ 0.05

Fail to reject H0. Flow rate does not significantly affect etch uniformity.

```

In [24]: # 4b
plt.figure(figsize=(7,5))
sns.boxplot(x="Flow", y="Uniformity", data=df)
plt.title("Boxplot of Etch Uniformity by Flow Rate")
plt.ylabel("Etch Uniformity (%)")
plt.xlabel("C2F6 Flow Rate")
plt.show()

```



```

In [25]: residuals = model.resid

```

```

In [26]: shapiro_test = stats.shapiro(residuals)
print("\nShapiro-Wilk Test for Normality:")
print(f"W = {shapiro_test.statistic:.4f}, p = {shapiro_test.pvalue:.4f}")
print(f"p = {shapiro_test.pvalue:.4f}")

```

Shapiro-Wilk Test for Normality:

W = 0.9605, p = 0.6111

p = 0.6111

```

In [27]: if shapiro_test.pvalue > 0.05:
    print(" Conclusion: Fail to reject H0: Normality assumption seems reasonable.")
else:
    print(" Conclusion: Reject H0: Non-normality in residuals.")

```

Conclusion: Fail to reject H0: Normality assumption seems reasonable.

```

In [28]: # 5a
strength = np.array([160,171,175,182,184,181,188,193,195,200])
pct_hardwood = np.array([10,15,15,20,20,20,25,25,28,30])

```

```

In [29]: data = pd.DataFrame({
    "Strength": strength,
    "PctHardwood": pct_hardwood
})

```

```

In [30]: print(data)
print("\n")

```

	Strength	PctHardwood
0	160	10
1	171	15
2	175	15
3	182	20
4	184	20
5	181	20
6	188	25
7	193	25
8	195	28
9	200	30

```
In [31]: X = sm.add_constant(pct_hardwood)
model = sm.OLS(strength, X).fit()
print(model.summary())
print("\n")
```

```

=====
                        OLS Regression Results
=====
Dep. Variable:          y      R-squared:          0.970
Model:                  OLS    Adj. R-squared:      0.966
Method:                 Least Squares    F-statistic:      260.0
Date:                  Thu, 25 Sep 2025    Prob (F-statistic):  2.20e-07
Time:                  07:06:40    Log-Likelihood:     -20.973
No. Observations:      10    AIC:              45.95
Df Residuals:          8    BIC:              46.55
Df Model:              1
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	143.8244	2.522	57.039	0.000	138.010	149.639
x1	1.8786	0.117	16.125	0.000	1.610	2.147

```

=====
Omnibus:                 1.211    Durbin-Watson:      2.132
Prob(Omnibus):           0.546    Jarque-Bera (JB):    0.701
Skew:                    0.157    Prob(JB):            0.704
Kurtosis:                1.742    Cond. No.            78.5
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

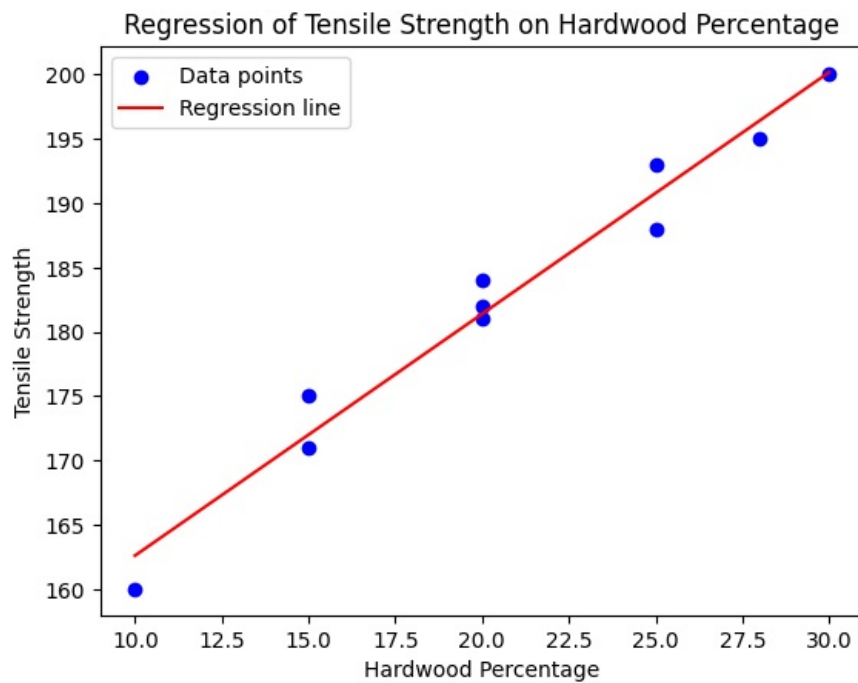
```
/opt/anaconda3/lib/python3.12/site-packages/scipy/stats/_axis_nan_policy.py:531: UserWarning: kurtosistest only
valid for n>=20 ... continuing anyway, n=10
res = hypotest_fun_out(*samples, **kws)
```

```
In [32]: # 5b
p_value = model.pvalues[1]
print(f"p-value for slope (PctHardwood): {p_value:.4f}")
if p_value < 0.05:
    print("Result: Reject H0. The slope is significant.")
else:
    print("Result: Fail to reject H0. No significant relationship detected.")
print("\n")
```

p-value for slope (PctHardwood): 0.0000

Result: Reject H0. The slope is significant.

```
In [33]: # 5c
plt.scatter(pct_hardwood, strength, color="blue", label="Data points")
plt.plot(pct_hardwood, model.predict(X), color="red", label="Regression line")
plt.xlabel("Hardwood Percentage")
plt.ylabel("Tensile Strength")
plt.title("Regression of Tensile Strength on Hardwood Percentage")
plt.legend()
plt.show()
```



```
In [34]: # 5d
print(f"R-squared: {model.rsquared:.3f}")
print(f"Estimated regression equation: Strength = {model.params[0]:.2f} + {model.params[1]:.2f} * PctHardwood")

if p_value < 0.05:
    print("There is a statistically significant positive relationship between hardwood percentage and strength.")
    print("The plot shows that as hardwood % increases, tensile strength tends to increase practically linearly")
else:
    print("No significant relationship between hardwood % and tensile strength was detected.")
```

R-squared: 0.970

Estimated regression equation: Strength = 143.82 + 1.88 * PctHardwood

There is a statistically significant positive relationship between hardwood percentage and strength.

The plot shows that as hardwood % increases, tensile strength tends to increase practically linearly.

In []:

In []:

In []:

In []: